

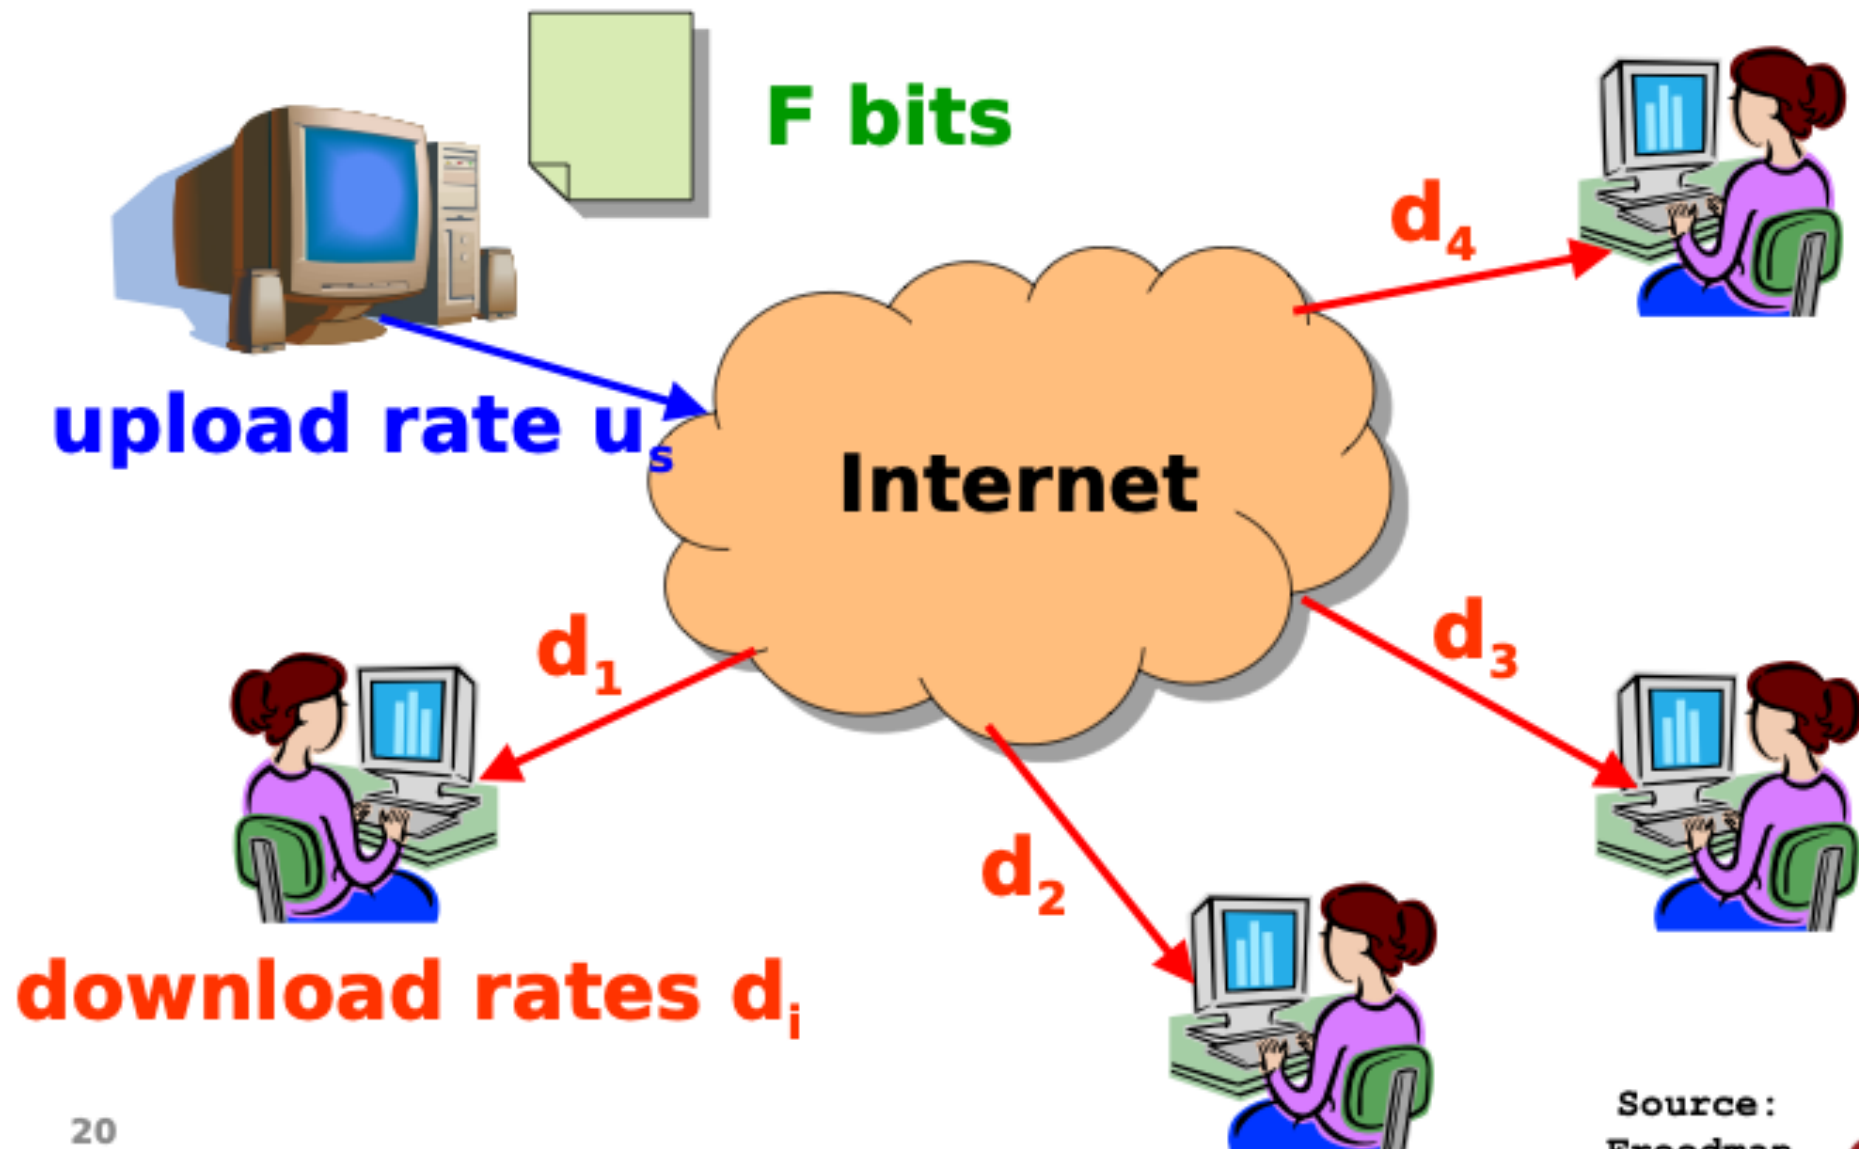


Application Layer: Peer to Peer File Sharing

Michael Kirkedal Thomsen

Based on slides compiled by Marcos Vaz Salles, adaptations by Vivek Shah and Michael Kirkedal Thomsen

Server Distributing a Large File



Server Distributing a Large File

- Sending an F -bit file to N receivers
 - Transmitting NF bits at rate u_s
 - ... takes at least NF/u_s time
- Receiving the data at the slowest receiver
 - Slowest receiver has download rate $d_{min} = \min_i \{d_i\}$
 - ... takes at least F/d_{min} time
- Download time: $\max\{NF/u_s, F/d_{min}\}$

Source:
Freedman



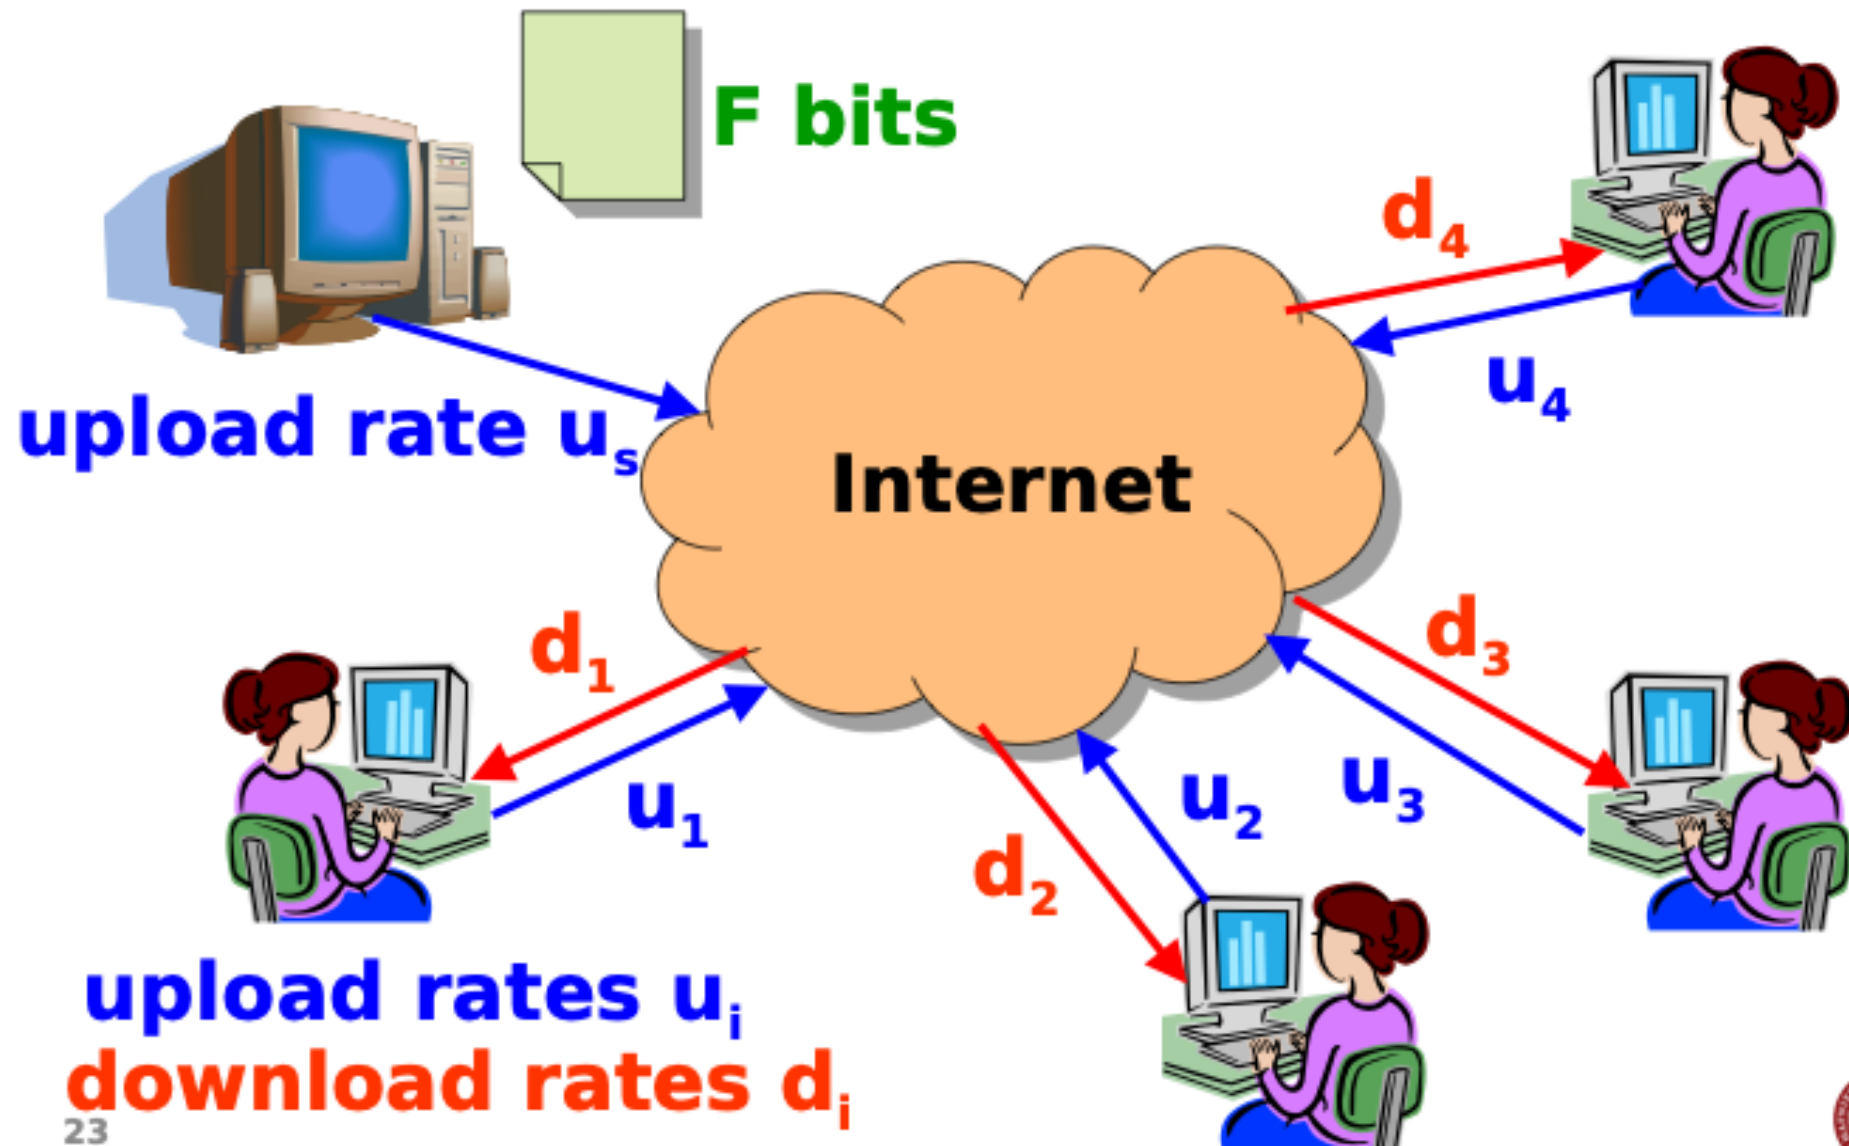
Speeding Up the File Distribution

- Increase the server upload rate
 - Higher link bandwidth at the server
 - Multiple servers, each with their own link
- Alternative: have the receivers help
 - Receivers get a copy of the data
 - ... and redistribute to other receivers
 - To reduce the burden on the server

Source:
Freedman



Pees Help Distributing a Large File



Peers Help Distributing a Large File

- Components of distribution latency
 - Server must send each bit: min time F/u_s
 - Slowest peer must receive each bit: min time F/d_{min}
- Upload time using all upload resources
 - Total number of bits: NF
 - Total upload bandwidth $u_s + \sum_i(u_i)$
 - Total: $\max\{F/u_s, F/d_{min}, NF/(u_s + \sum_i(u_i))\}$
- *Peer to peer is self-scaling*
 - *Download time grows slowly with N*
 - Client-server: $\max\{NF/u_s, F/d_{min}\}$
 - Peer-to-peer: $\max\{F/u_s, F/d_{min}, NF/(u_s + \sum_i(u_i))\}$



Peer-to-Peer Networks: BitTorrent

- BitTorrent history
 - 2002: B. Cohen debuted BitTorrent
- Emphasis on efficient fetching, not searching
 - Distribute same file to many peers
 - Single publisher, many downloaders
- Preventing free-loading
 - Incentives for peers to contribute

Source:
Freedman



BitTorrent: Tracker

- Infrastructure node
 - Keeps track of peers participating in the torrent
 - Peers register with the tracker when it arrives
- Tracker selects peers for downloading
 - Returns a random set of peer IP addresses
 - So the new peer knows who to contact for data
- Can have “trackerless” system
- Using distributed hash tables (DHTs)

Source:
Freedman



BitTorrent: Chunk Request Order

- Which chunks to request?
 - Could download in order
 - Like an HTTP client does
- Problem: many peers have the early chunks
 - Peers have little to share with each other
 - Limiting the scalability of the system
- Problem: eventually nobody has rare chunks
 - E.g., the chunks need the end of the file
 - Limiting the ability to complete a download
- Solutions: random selection and rarest first

Source: Freedman



BitTorrent: Rarest Chunk First

- Which chunks to request first?
 - Chunk with fewest available copies (i.e., rarest chunk)
- Benefits to the peer
 - Avoid starvation when some peers depart
- Benefits to the system
 - Avoid starvation across all peers wanting a file
 - Balance load by equalizing # of copies of chunks

Source: Freedman



Free-Riding in P2P Networks

- Vast majority of users are free-riders
 - Most share no files and answer no queries
 - Others limit # of connections or upload speed
- A few “peers” essentially act as servers
 - A few individuals contributing to the public good
 - Making them hubs that basically act as a server
- BitTorrent prevent free riding
 - Allow the fastest peers to download from you
 - Occasionally let some free loaders download

Source: Freedman



Bit-Torrent: Preventing Free-Riding

- Peer has limited upload bandwidth
 - And must share it among multiple peers
 - Tit-for-tat: favor neighbors uploading at highest rate
- Rewarding the top four neighbors
 - Measure download bit rates from each neighbor
 - Reciprocate by sending to the top four peers
- Optimistic unchoking
 - Randomly try a new neighbor every 30 seconds
 - So new neighbor has a chance to be a better partner
 - Compatible peers find each other

Source: Freedman



Peer-to-Peer Naming

- But...
 - Peers may come and go
 - Peers need to find each other
 - Peers need to be willing to help each other

Source: Freedman



Locating the Relevant Peers

- Three main approaches
 - Central directory (Napster)
 - Query flooding (Gnutella)
 - Hierarchical overlay (Kazaa, modern Gnutella)
- Design goals
 - Scalability
 - Simplicity
 - Robustness
 - Plausible deniability

Source: Freedman



Recap: Streaming multimedia: DASH

- *DASH*: *D*ynamic, *A*daptive *S*treaming over *H*TTP
- *server*:
 - divides video file into multiple chunks
 - each chunk stored, encoded at different rates
- *manifest file*:
 - provides URLs for different chunks
- *client*:
 - periodically measures server-to-client bandwidth
 - consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)

Source: Kurose & Ross



P2P media streaming: Octoshape

- Commercial hosting not cable to home, founded 2003
 - Co-founder Stephen Alstrup
 - Broadcasting fee is paid by broadcasters
 - Free for consumers
- Audio and Video, 32kbps to 800kbps
- Mesh based, bit-torrent like, Content Server pushes content to some peers, it propagates from there
- Peers advertise their joining to everyone (?)
 - Probably to some network or geographic topography, esp. if the live content is popular, this can cause a notification storm.
[speculation]
- Eurovision in 2006, 2007, Tour de France in 2007 with ~1.5Mbps High Quality
- Requires a web browser, any streaming client and Octoshape client

Source: Jörg Ott

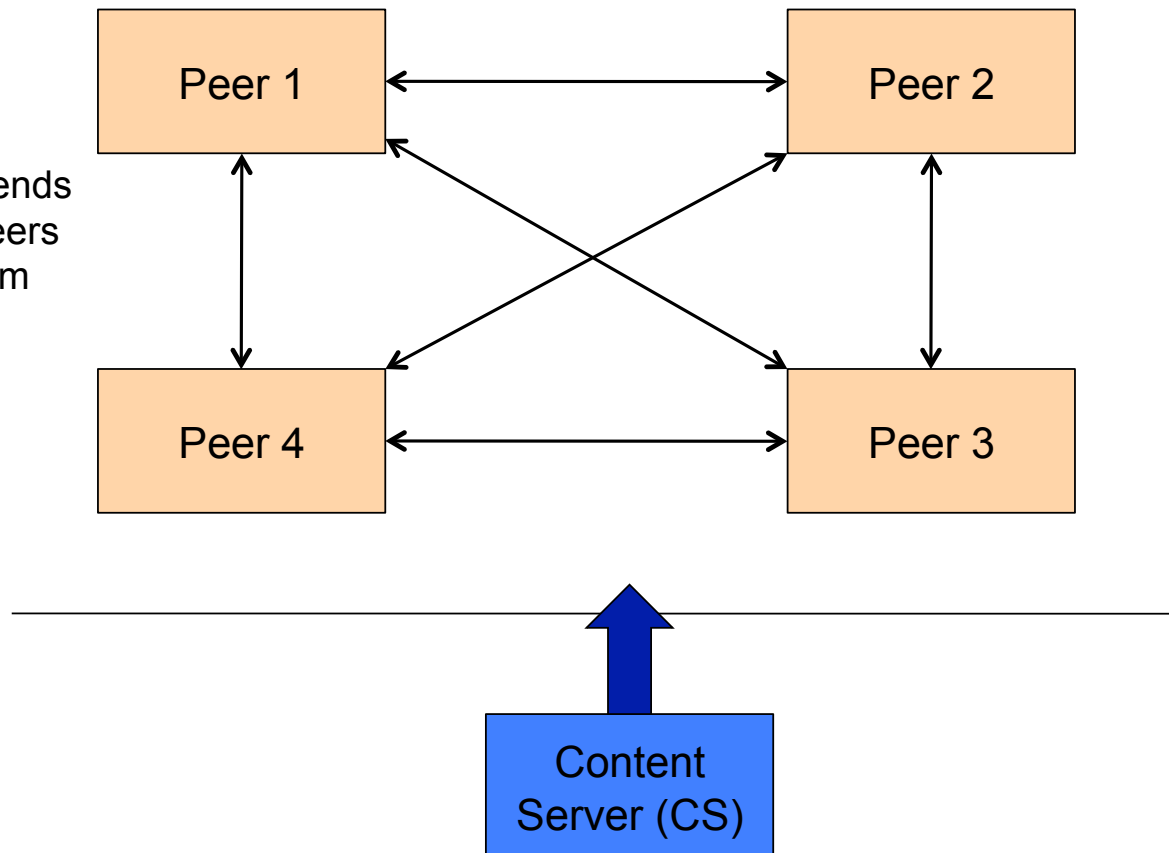


P2P media streaming: Octoshape

Peer decides to split the streams in to k-parts.

- 1) Stand-by peers
- 2) Connected peers

No of parts depends on number of peers downloading from that peer



Source: Jörg Ott



Summary

- P2P applications
 - Self scalability
 - BitTorrent – Popular P2P file sharing protocol
 - Rarest chunk first, fair trading + optimistic unchoking
 - Many implementation of media streaming
 - Today hybrids via Amazon services

